Baseline K = 20, depth (8,2), 5 qubits, trained weights (iteration 10,000)
581 gates, 192 two-qubit gates, depth 355, two-qubit depth 178. q0 = color qubit (target of every CRY), q1-q4 = address qubits. Measured in order q1 q2 q3 q4 q0. RX = noise input, RY and CRY = trained weights.
Used in: 4x4 baseline training, Aer and fake_torino export runs.

